

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

# Constructing a Knowledge Base on Aging

## An Automated Approach

Mark Farrell

Bioinformatics Researcher

Center for Research and Education on Aging  
Lawrence Berkeley National Laboratory  
University of California, Berkeley

September 4th, 2014

# Outline

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## 1 About Me

## 2 Automatically Constructing Knowledge Bases

## 3 Extracting Facts in a Structured Format

## 4 Results & Discussion

# Who I am

## Mark Farrell

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion



# Where I'm from

## Atlantic Canada

Constructing a  
Knowledge  
Base on Aging

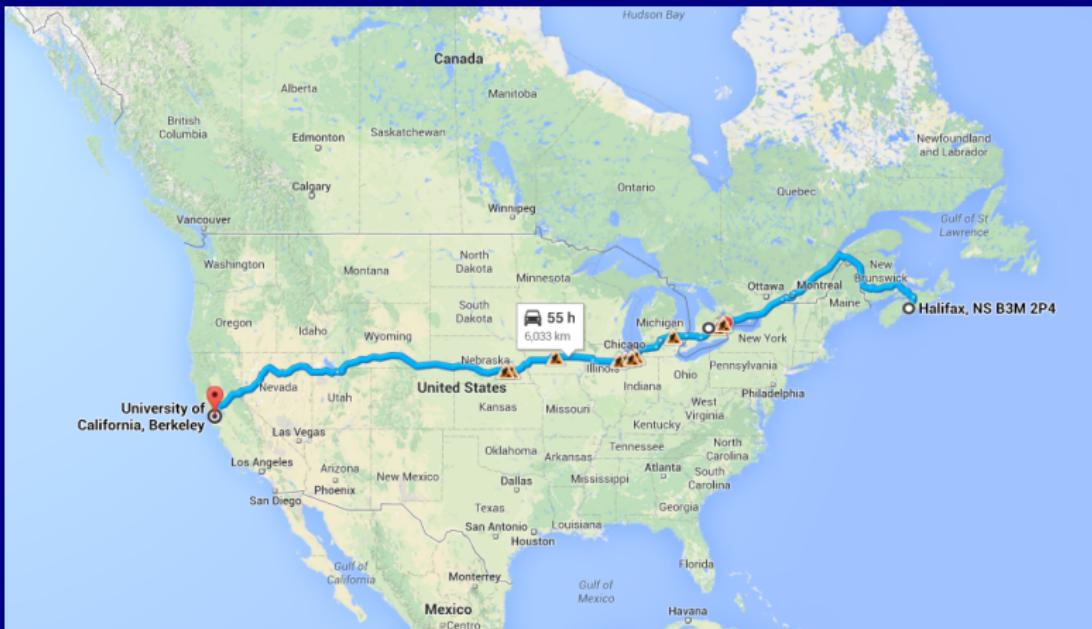
Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion



# What I'm Studying

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Year 2

Program Bachelor of Computer Science

Faculty Mathematics

Institution University of Waterloo

Location Waterloo, Ontario, Canada



# How I Became a Researcher at CREA

Prof. Garan Emailed the University of Waterloo Computer Science Club

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion



# Why was I Interested in Research at CREA?

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Hobby	CREA Research
Reverse Engineering	Natural Language Processing
Game Development	Artificial Intelligence
Health	Bioinformatics

# Outline

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## 1 About Me

## 2 Automatically Constructing Knowledge Bases

## 3 Extracting Facts in a Structured Format

## 4 Results & Discussion

# Overview

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- CREA is constructing a knowledge base to study and understand the human aging process.

# Overview

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- CREA is constructing a knowledge base to study and understand the human aging process.
- New discoveries are published quickly and in large volume.

# Overview

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- CREA is constructing a knowledge base to study and understand the human aging process.
- New discoveries are published quickly and in large volume.
- It is infeasible to construct the knowledge base by hand.

# Overview

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- CREA is constructing a knowledge base to study and understand the human aging process.
- New discoveries are published quickly and in large volume.
- It is infeasible to construct the knowledge base by hand.
- Working on software to construct the knowledge base automatically.

# Introduction

## How to Automatically Construct the Knowledge Base

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Routinely search for keywords related to aging, dowloading text articles from sources like PubMed.

# Introduction

## How to Automatically Construct the Knowledge Base

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Routinely search for keywords related to aging, dowloading text articles from sources like PubMed.
- Build a spam filter to get rid of non-scientific sentences.

# Introduction

## How to Automatically Construct the Knowledge Base

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Routinely search for keywords related to aging, dowloading text articles from sources like PubMed.
- Build a spam filter to get rid of non-scientific sentences.
- Extract scientific facts from the sentences and save them in a structured format.

# Introduction

## How to Automatically Construct the Knowledge Base

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Routinely search for keywords related to aging, dowloading text articles from sources like PubMed.
- Build a spam filter to get rid of non-scientific sentences.
- Extract scientific facts from the sentences and save them in a structured format.
- Provide a graphical interface that allows users to search and otherwise explore the knowledge base.

# Summary of Progress

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Devised and implemented the method for finding simple facts in sentences, extracting them in a structured format.

# Summary of Progress

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Devised and implemented the method for finding simple facts in sentences, extracting them in a structured format.
- Began work on a web viewer for the knowledge base.

# Outline

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## 1 About Me

## 2 Automatically Constructing Knowledge Bases

## 3 Extracting Facts in a Structured Format

## 4 Results & Discussion

# Extracting Facts in a Structured Format

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Tokenization   Input a text document and read it, one sentence  
at a time.

# Extracting Facts in a Structured Format

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

**Tokenization** Input a text document and read it, one sentence at a time.

**Parsing** For each sentence, generate a constituent tree that describes its phrase structure.

# Extracting Facts in a Structured Format

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Tokenization      Input a text document and read it, one sentence at a time.

Parsing      For each sentence, generate a constituent tree that describes its phrase structure.

Compilation      Extract facts by pattern matching on each constituent tree.

# Constituent Tree Tags

## Word Level

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Tag	Meaning	Example
DT	Determiner	the
IN	Preposition	of
JJ	Adjective	blue
RB	Adverb	quickly
CC	Coordinating Conjunction	and
NN	Singular Noun	monkey
NNS	Plural Noun	monkeys
VB	Base Verb	fall
VBD	Singular Present Verb	falls
VBN	Past Tense Verb	fell
VBG	Past Participle Verb	fallen
...	Gerund Verb	falling
...	...	...

# Constituent Tree Tags

## Phrase Level

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Tag	Meaning	Example
NP	Noun Phrase	the woman
VP	Verb Phrase	calls the man
PP	Prepositional Phrase	from the store
ADVP	Adverb Phrase	quickly and quietly
ADJP	Adjective Phrase	blue and red
CONJP	Conjunctive Phrase	as well as
...	...	...

# Constituent Tree Tags

## Clause Level

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

Tag	Meaning	Example
S	Declarative Clause	the dog walks
SBAR	Conjunction + Clause	that the dog walks
...	...	...

# Information Extraction

## Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

### Input Token

The man walks the dog.

# Information Extraction

## Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

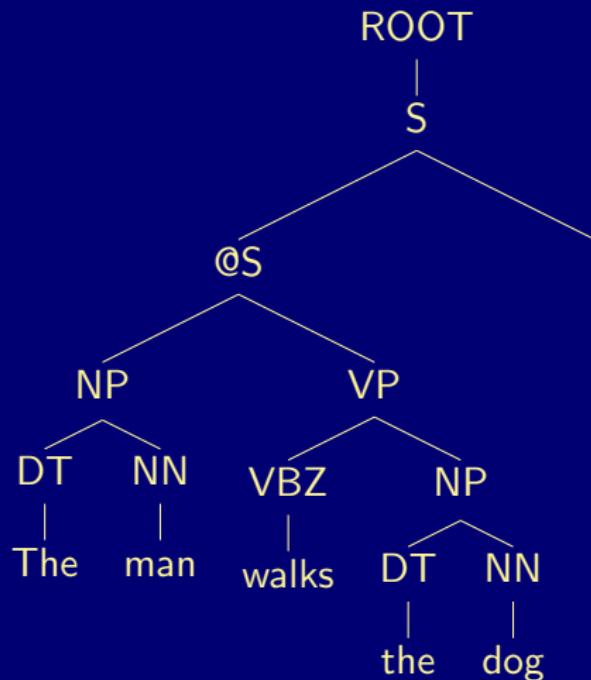
About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

### Parse Token



# Information Extraction

## Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

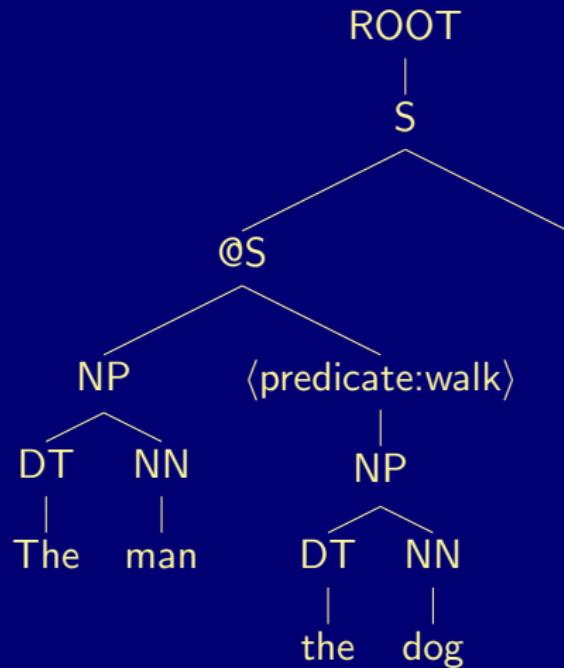
About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## Compile Token



# Information Extraction

## Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

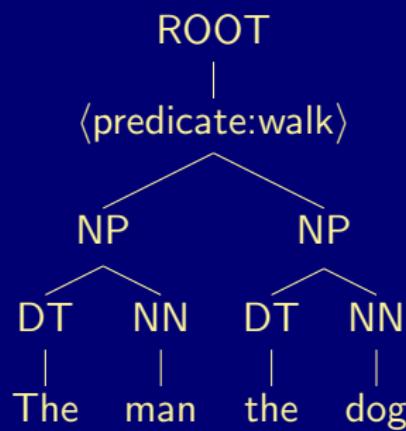
About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

### Compile Token



# Information Extraction

## Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

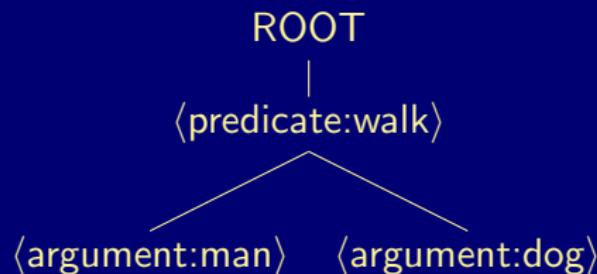
About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

### Compile Token



# Information Extraction Example

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## Output Fact

`walk(man, dog).`

# Outline

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

## 1 About Me

## 2 Automatically Constructing Knowledge Bases

## 3 Extracting Facts in a Structured Format

## 4 Results & Discussion

# Software Demonstration

A preview of CREA's knowledge base, compiled from PubMed abstracts.

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion



# Performance

## Parallelization

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- It is possible to extract facts from many sentences at the same time.

# Accuracy

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Filter spam sentences from documents.

# Accuracy

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Filter spam sentences from documents.
- The accuracy of the parser could be optimized:

# Accuracy

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Filter spam sentences from documents.
- The accuracy of the parser could be optimized:
  - Should be trained to identify more nouns from the biomedical domain.

# Accuracy

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Filter spam sentences from documents.
- The accuracy of the parser could be optimized:
  - Should be trained to identify more nouns from the biomedical domain.
- Define more patterns for extracting facts:

# Accuracy

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Filter spam sentences from documents.
- The accuracy of the parser could be optimized:
  - Should be trained to identify more nouns from the biomedical domain.
- Define more patterns for extracting facts:
  - The software succeeds around 50% of the time.

# Missing Features

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Support negated clauses and conditional logic.

# Missing Features

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Support negated clauses and conditional logic.
- Facts can contradict each other:

# Missing Features

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Support negated clauses and conditional logic.
- Facts can contradict each other:
  - Store the probability that is true as the weight of its edge on the knowledge base's graph.

# Missing Features

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Support negated clauses and conditional logic.
- Facts can contradict each other:
  - Store the probability that is true as the weight of its edge on the knowledge base's graph.
- Scale and launch the software service.

# Conclusion

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Demonstrated a method for automatically constructing CREA's knowledge base on aging.

# Conclusion

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Demonstrated a method for automatically constructing CREA's knowledge base on aging.
- Showed how to extract facts from English text in the knowledge base's structured format.

# Conclusion

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Demonstrated a method for automatically constructing CREA's knowledge base on aging.
- Showed how to extract facts from English text in the knowledge base's structured format.
- Discussed the need to improve software accuracy by lensing in on the biomedical domain.

# Conclusion

Constructing a  
Knowledge  
Base on Aging

Mark Farrell

About Me

Automatically  
Constructing  
Knowledge  
Bases

Extracting  
Facts in a  
Structured  
Format

Results &  
Discussion

- Demonstrated a method for automatically constructing CREA's knowledge base on aging.
- Showed how to extract facts from English text in the knowledge base's structured format.
- Discussed the need to improve software accuracy by lensing in on the biomedical domain.
- Suggested how the software implementation can be scaled for production usage.